

# Project2: Static routing forwarding

---

## Project2: Static routing forwarding

一 实验目的

二 实验环境

三 实验内容

四 实验流程

实验待完善内容为：

4.1 相关代码的TODO部分

五 实验结果

5.1 实验一

5.2 实验二

徐磊 201828018670040

崔天宇 201818018670007

陈洁 2018E8018661111

## 一 实验目的

---

- 理解 IP、ARP、ICMP 协议的工作机制
- 实现 IP 地址查找与数据包转发
- 实现 ARP 请求和应答、ARP 数据包缓存管理
- 实现 ICMP 消息的发送
- 学习 ping, traceroute 等网络命令的基本功能和使用方法

## 二 实验环境

---

- Ubuntu16.04
- Mininet

## 三 实验内容


---

根据已有的基础代码，继续完成 TODO 部分。在两种网络拓扑结构下，主要完成以下三个部分：

- **IP协议相关：**查找转发(基于最长前缀匹配规则)
- **ARP协议相关：**IP-Mac 地址映射基本概念，ARP 请求、ARP 回应数据包格式，ARP 条目查询，ARP 缓存管理

- **ICMP 协议**：实验中涉及到的几种 ICMP 数据包格式(路由表查找失败、ARP 查 询失败、TTL 值减为 0、收到 ping 本端又的数据包)

实验一：对应的星型网络拓扑结构：



完善代码中的 TODO 部分，实现路由器的连通性测试。即在 h1 上进行 ping 实验：

Ping 10.0.1.1 (r1)，可以 ping 通

Ping 10.0.2.22 (h2)，可以 ping 通

Ping 10.0.3.33 (h3)，可以 ping 通

Ping 10.0.3.11，返回 ICMP Destination Host Unreachable

Ping 10.0.4.1，返回 ICMP Destination Net Unreachable

实验二:根据要求配置网络拓扑环境:



建立如上线性网络拓扑，实现路由器的连通性测试。以及路径测试:在一台主 机上可以 traceroute 其他主机，获取数据包经过每个节点入端又的 IP 地址。

## 四 实验流程

实验待完善内容为：

- **处理IP数据包**：

判断是否为 ICMP echo 请求，否则利用最长前缀匹配规则转发 IP 包，转发 IP 包 之前需要检查 TTL，更新 checksum 等

- **发送 ICMP 数据包**：

满足下述四个条件时发送 ICMP 数据包

- TTL 值为 0
- 查找不到路由表条目收到
- ping 本端又的包
- ARP 查询失败

- **ARP 缓存管理:**

在没有触发前三个条件时，将数据包添加到 ARP 缓存的等待队列中，并发送 ARP 请求。注意缓存管理要避免死锁的出现(lock unlock)

- **处理 ARP 请求和应答:**

收到 ARP 数据包时，先判断是 ARP 请求还是 ARP 应答

处理 ARP 请求数据包时，先发送 ARP 回复数据包，再将 IP-Mac 地址映射对插入 ARP 缓存中，并查找等待队列中是否有数据包符合地址映射对，如有，则将对的数据包发送 处理 ARP 回复数据包时，同理也先将 IP-Mac 地址映射对插入 ARP 缓存中，再查找等待队列中是否有数据包符合地址映射对，如有，则将对的数据包发送。

## 4.1 相关代码的TODO部分

代码：arp.c

- 发送arp请求

```
// - send an arp request: encapsulate an arp request packet, send it out through
// - iface_send_packet
void arp_send_request(iface_info_t *iface, u32 dst_ip)
{
    fprintf(stderr, "TODO: send arp request when lookup failed in arpcache.\n");
    char* packet;
    packet = (char*)malloc(sizeof(struct ether_header) + sizeof(struct ether_arp));
    struct ether_header* eh;
    eh = (struct ether_header*)(packet);
    memcpy(eh->ether_shost, iface->mac, ETH_ALEN*sizeof(u8));
    show(eh->ether_shost);
    u8 temp = 255;
    for(int q = 0; q < ETH_ALEN; q++) eh->ether_dhost[q] = temp;
    show(eh->ether_dhost);
    eh->ether_type = htons(ETH_P_ARP);

    struct ether_arp* eh_arp;
    eh_arp = (struct ether_arp*)(packet + ETHER_HDR_SIZE);
    eh_arp->arp_hrd = htons(0x01);
    eh_arp->arp_pro = htons(0x0800);
    eh_arp->arp_hln = 6;
    eh_arp->arp_pln = 4;
    eh_arp->arp_op = htons(0x01);
    memcpy(eh_arp->arp_sha, iface->mac, ETH_ALEN*sizeof(u8));
    printf("*****\n");
    show(iface->mac);
    show(eh_arp->arp_sha);
    eh_arp->arp_spa = htonl(iface->ip);
    eh_arp->arp_tpa = htonl(dst_ip);
    memset(eh_arp->arp_tha, 0, ETH_ALEN*sizeof(u8));
    // show(eh_arp->arp_tha);
    iface_send_packet(iface, packet, ETHER_HDR_SIZE+sizeof(struct ether_arp));
}
```

- 发送arp reply

```

// send an arp reply packet: encapsulate an arp reply packet, send it out
// through iface_send_packet
void arp_send_reply(iface_info_t *iface, struct ether_arp *req_hdr)
{
    fprintf(stderr, "T000: send arp reply when receiving arp request.\n");
    char * packet;
    packet = (char *) malloc(ETHER_HDR_SIZE + sizeof(struct ether_arp));
    struct ether_header * eh;
    eh = (struct ether_header *) packet;
    struct ether_arp * eh_arp;
    eh_arp = (struct ether_arp *) (packet + ETHER_HDR_SIZE);
    memcpy(eh->ether_shost, iface->mac, ETH_ALLEN * sizeof(u8));
    memcpy(eh->ether_dhost, req_hdr->arp_sha, ETH_ALLEN * sizeof(u8));
    eh->ether_type = htons(ETH_P_ARP);
    eh_arp->arp_hrd = htons(0x01);
    eh_arp->arp_pro = htons(0x0800);
    eh_arp->arp_hln = 6;
    eh_arp->arp_pln = 4;
    eh_arp->arp_op = htons(ARPOP_REPLY);
    eh_arp->arp_spa = htonl(iface->ip);
    u32 my_ip = ntohl(req_hdr->arp_spa);
    eh_arp->arp_tpa = htonl(my_ip);
    memcpy(eh_arp->arp_sha, iface->mac, ETH_ALLEN * sizeof(u8));
    memcpy(eh_arp->arp_tha, req_hdr->arp_sha, ETH_ALLEN * sizeof(u8));

    iface_send_packet(iface, packet, sizeof(struct ether_arp) + ETHER_HDR_SIZE);
}

```

- 发送arp request和arp reply，并且将ip->mac映射和相关信息放入arpcache中

```

void handle_arp_packet(iface_info_t *iface, char *packet, int len)
{
    fprintf(stderr, "T000: process arp packet: arp request & arp reply.\n");
    struct ether_header * eh;
    struct ether_arp * eh_arp;
    eh = (struct ether_header *) (packet);
    eh_arp = (struct ether_arp *) (packet + ETHER_HDR_SIZE);
    if (ntohs(eh_arp->arp_op) == 0x01) {
        if (ntohl(eh_arp->arp_tpa) == iface->ip) {
            arp_send_reply(iface, eh_arp);
            arpcache_insert(ntohl(eh_arp->arp_spa), eh_arp->arp_sha);
        }
    }
    else if (ntohs(eh_arp->arp_op) == 0x02) {
        if (ntohl(eh_arp->arp_tpa) == iface->ip) {
            arpcache_insert(ntohl(eh_arp->arp_spa), eh_arp->arp_sha);
        }
    }
    free(packet);
}

```

```

// send (IP) packet through arpcache lookup
//
// Lookup the mac address of dst_ip in arpcache. If it is found, fill the
// ethernet header and emit the packet by iface_send_packet, otherwise, pending
// this packet into arpcache, and send arp request.
void iface_send_packet_by_arp(iface_info_t *iface, u32 dst_ip, char *packet, int len)
{
    struct ether_header * eh = (struct ether_header *) packet;
    memcpy(eh->ether_shost, iface->mac, ETH_ALLEN);
    eh->ether_type = htons(ETH_P_IP);
    u8 dst_mac[ETH_ALLEN];
    int found = arpcache_lookup(dst_ip, dst_mac);
    if (!found) {
        memcpy(eh->ether_dhost, dst_mac, ETH_ALLEN * sizeof(u8));
        if (ntohs(eh->ether_type) == ETH_P_IP) {
            struct iphdr * myiph;
            myiph = (struct iphdr *) (packet + ETHER_HDR_SIZE);
            struct icmp_hdr * myicmph;
            myicmph = (struct icmp_hdr *) (packet + ETHER_HDR_SIZE + IP_HDR_SIZE(myiph));
            u32 sip = ntohl(myiph->saddr);
            u32 dip = ntohl(myiph->daddr);
        }
        iface_send_packet(iface, packet, len);
    }
    else {
        arpcache_append_packet(iface, dst_ip, packet, len);
    }
}

```

代码 Arpcache.c

- 查找arpchae中已将有相同的IP项

```

// lookup the IP->mac mapping
//
// traverse the table to find whether there is an entry with the same IP
// and mac address with the given arguments
int arpcache_lookup(u32 ip4, u8 mac[ETH_ALEN])
{
    pthread_mutex_lock(&arpcache.lock);
    fprintf(stderr, "TODO: lookup ip address in arp cache.\n");
    show_cache(arpcache);
    int i=0;
    for(i=0; i < MAX_ARP_SIZE; i++){
        IP_int2string(ip4);
        if(arpcache.entries[i].ip4 == ip4 && arpcache.entries[i].valid == 1){
            printf("find in cache!\n");
            // pthread_mutex_lock(&arpcache.lock);
            memcpy(mac, arpcache.entries[i].mac, ETH_ALEN*sizeof(u8));
            // printf("-----\n");
            // show_mac(mac);
            // show_mac(arpcache.entries[i].mac);
            pthread_mutex_unlock(&arpcache.lock);
            return 1;
        }
    }
    printf("did not find in cache\n");
    pthread_mutex_unlock(&arpcache.lock);
    return 0;
}

```

- Append the packet to arpcache

```

// append the packet to arpcache
//
// Lookup in the list which stores pending packets, if there is already an
// entry with the same IP address and iface (which means the corresponding arp
// request has been sent out), just append this packet at the tail of that entry
// (the entry may contain more than one packet); otherwise, malloc a new entry
// with the given IP address and iface, append the packet, and send arp request.
void arpcache_append_packet(iface_info_t *iface, u32 ip4, char *packet, int len)
{
    fprintf(stderr, "TODO: append the ip address if lookup failed, and send arp request if necessary.\n");
    // find
    pthread_mutex_lock(&arpcache.lock);
    struct arp_req *req_entry = NULL, *req_q;
    list_for_each_entry_safe(req_entry, req_q, &(arpcache.req_list), list){
        if(req_entry->iface->ip == iface->ip && req_entry->ip4 == ip4){
            struct cached_pkt *cache;
            cache = (struct cached_pkt *)malloc(sizeof(struct cached_pkt));
            cache->packet = packet;
            cache->len = len;
            list_add_tail(&(cache->list), &(req_entry->cached_packets));
            pthread_mutex_unlock(&arpcache.lock);
            return;
        }
    }
    // add
    struct arp_req req_q_new;
    req_q_new = (struct arp_req)malloc(sizeof(struct arp_req));
    req_q_new->iface = iface;
    req_q_new->ip4 = ip4;
    req_q_new->sent = time(NULL);
    req_q_new->retries = 0;
    init_list_head(&(req_q_new->cached_packets));
    list_add_tail(&(req_q_new->list), &(arpcache.req_list));

    struct cached_pkt *cache_new;
    cache_new = (struct cached_pkt *)malloc(sizeof(struct cached_pkt));
    cache_new->packet = packet;
    cache_new->len = len;
    list_add_tail(&(cache_new->list), &(req_q_new->cached_packets));
    pthread_mutex_unlock(&arpcache.lock);
    arp_send_request(iface, ip4);
}

```

将IP->mac的映射插入arpcache中，如果arpcache中没有位置可以添加映射，则随机插入某项。如果有待处理的数据包等待这个映射，请为每个数据填充以太网头，并将它们发送出去

```

//insert the IP->mac mapping into arpcache, if there are pending packets
//waiting for this mapping, fill the ethernet header for each of them, and send
//them out
void arpcache_insert(u32 ip4, u8 mac[ETH_ALEN])
{
    fprintf(stderr, "T000: insert ip->mac entry, and send all the pending packets.\n");
    //pthread_mutex_lock(&arpcache.lock);
    //insert the IP->mac mapping into arpcache
    show_cache(arpcache);
    int flag = 0;
    for(int i = 0; i < MAX_ARP_SIZE; i++){
        pthread_mutex_lock(&arpcache.lock);
        if(arpcache.entries[i].ip4 == 0){
            pthread_mutex_lock(&arpcache.lock);
            if(isRepeat(ip4))
            {
                arpcache.entries[i].ip4 = ip4;
                memcpy(arpcache.entries[i].mac, mac, ETH_ALEN*sizeof(u8));
                arpcache.entries[i].added = time(NULL);
                arpcache.entries[i].valid = 1;
            }
            pthread_mutex_unlock(&arpcache.lock);
            flag = 1;
            break;
        }
        else
            pthread_mutex_unlock(&arpcache.lock);
    }
    //arpcache has been filled totally, insert randomly
    if(!flag){
        int index = rand()%MAX_ARP_SIZE;
        pthread_mutex_lock(&arpcache.lock);
        arpcache.entries[index].ip4 = ip4;
        memcpy(arpcache.entries[index].mac, mac, ETH_ALEN*sizeof(u8));
        arpcache.entries[index].added = time(NULL);
        arpcache.entries[index].valid = 1;
        pthread_mutex_unlock(&arpcache.lock);
    }
}

//deal with pending packets
struct arp_req *req_entry = NULL, *req_q;
list_for_each_entry_safe(req_entry, req_q, &(arpcache.req_list), list){
    if(req_entry->ip4 == ip4){
        struct cached_pkt *cache = NULL, *cache_q;
        list_for_each_entry_safe(cache, cache_q, &(req_entry->cached_packets), list){
            char *packet = cache->packet;
            int len = cache->len;
            struct ether_header *eh;
            eh = (struct ether_header *) (packet);
            memcpy(eh->ether_dhost, mac, ETH_ALEN*sizeof(u8));
            iface_send_packet(req_entry->iface, packet, len);
            list_delete_entry(&(cache->list));
        }
        list_delete_entry(&(req_entry->list));
    }
}
//show_cache(arpcache);
}

```

- 定期扫描arpcache

对于IP-> mac条目，如果条目在表中已超过15秒，将其从表中删除。对于挂起的数据包，如果arp请求在1秒前发送出去，而未收到回复，则重新发送arp请求。如果arp请求已经发送了5次而没有收到arp回复，则对于每个挂起的数据包，发送icmp数据包（DEST\_HOST\_UNREACHABLE），并丢弃这些数据包。

```

// sweep arpcache periodically
//
// For the IP->mac entry, if the entry has been in the table for more than 15
// seconds, remove it from the table.
// For the pending packets, if the arp request is sent out 1 second ago, while
// the reply has not been received, retransmit the arp request. If the arp
// request has been sent 5 times without receiving arp reply, for each
// pending packet, send icmp packet (DEST_HOST_UNREACHABLE), and drop these
// packets.
void *arpcache_sweep(void *arg)
{
    while (1) {
        sleep(1);
        fprintf(stderr, "T000: sweep arpcache periodically: remove old entries, resend arp requests.\n");
        show_cache(arpcache);
        // 避免死锁
        pthread_mutex_unlock(&arpcache.lock);
        pthread_mutex_lock(&arpcache.lock);
        // if the entry has been in the table for more than 15
        // seconds, remove it from the table
        for (int i = 0; i < MAX_APP_SIZE; i++) {
            if (time(NULL) - arpcache.entries[i].added >= 15 && arpcache.entries[i].ip4 != 0) {
                arpcache.entries[i].ip4 = 0;
                arpcache.entries[i].valid = 0;
            }
        }
        // For the pending packets, if the arp request is sent out 1 second ago, while the reply has not been received,
        struct arp_req *arp_req, *arp_req_q;
        time_t nowtime = time((time_t *) NULL);
        list_for_each_entry_safe(arp_req, arp_req_q, &(arpcache.req_list), list) {
            if (nowtime - arp_req->sent > 1 && arp_req->retries < 5) {
                iface_info_t *iface = arp_req->iface;
                u32 ip4 = arp_req->ip4;
                arp_send_request(iface, ip4);
                arp_req->sent = nowtime;
                arp_req->retries += 1;
            }
        }
        // more than 5, resend arp requests
        else if (arp_req->retries > 2) {
            struct cached_pkt *cache, *cache_q;
            list_for_each_entry_safe(cache, cache_q, &(arp_req->cached_packets), list) {
                char *packet = (char *) cache->packet;
                struct ether_header *eh = (struct ether_header *) packet;
                struct iphdr *iph = packet_to_ip_hdr(packet);
                u32 source_ip = ntohl(iph->saddr);
                u32 dst_ip = ntohl(iph->daddr);
                pthread_mutex_unlock(&arpcache.lock);
                u32 dst = ntohl(iph->saddr);
                rt_entry_t *entry = longest_prefix_match(dst);
                u32 sip = entry->iface->ip;
                printf("sending ICMP packet\n");
                icmp_send_packet((char *) cache->packet, cache->len, 3, 1, sip);
                pthread_mutex_lock(&arpcache.lock);
                list_delete_entry(&(cache->list));
            }
            pthread_mutex_unlock(&arpcache.lock);
            pthread_mutex_lock(&arpcache.lock);
            list_delete_entry(&(arp_req->list));
        }
        pthread_mutex_unlock(&arpcache.lock);
    }
    return NULL;
}

```

代码 lcmp.c

- 发送ICMP包



```

// send icmp packet
void icmp_send_packet(const char *in_pkt, int len, u8 type, u8 code, u32 sip)
{
    fprintf(stderr, "T000: malloc and send icmp packet.\n");
    int offset;
    char *icmp_packet;
    struct ether_header *eh_temp = (struct ether_header *)in_pkt;
    struct iphdr *iph_temp = (struct iphdr *)in_pkt + ETHER_HDR_SIZE;
    struct icmp_hdr *icmp_temp = (struct icmp_hdr *)in_pkt + ETHER_HDR_SIZE + IP_HDR_SIZE(iph_temp);
    if (type == 0 && code == 0) {
        icmp_packet = (char *)malloc(len);
        struct ether_header *eh = (struct ether_header *)icmp_packet;
        struct iphdr *iph = (struct iphdr *)icmp_packet + ETHER_HDR_SIZE;
        struct icmp_hdr *icmph = (struct icmp_hdr *)icmp_packet + ETHER_HDR_SIZE + IP_HDR_SIZE(iph_temp);
        eh->ether_type = htons(ETH_P_IP);
        memcpy(eh->ether_shost, eh_temp->ether_shost, ETH_ALEN * sizeof(u8));
        memcpy(eh->ether_dhost, eh_temp->ether_dhost, ETH_ALEN * sizeof(u8));
        u32 daddr_temp = sip;
        u32 saddr_temp = ntohl(iph_temp->saddr);
        ip_init_hdr(iph, daddr_temp, saddr_temp, len - ETHER_HDR_SIZE, IPPROTO_ICMP);
        icmph->type = 0;
        icmph->code = 0;

        memcpy(((char *)icmph) + 4, ((char *)icmp_temp) + 4, len - ETHER_HDR_SIZE - IP_HDR_SIZE(iph_temp));

        icmph->checksum = icmp_checksum(icmph, len - ETHER_HDR_SIZE - IP_HDR_SIZE(iph_temp));
        offset = len;
    }
    else {
        offset = ETHER_HDR_SIZE + 2 * IP_HDR_SIZE(iph_temp) + 16;
        icmp_packet = (char *)malloc(ETHER_HDR_SIZE + IP_HDR_SIZE(iph_temp) + 8 + IP_HDR_SIZE(iph_temp) + 8);
        struct ether_header *eh = (struct ether_header *)icmp_packet;
        struct iphdr *iph = (struct iphdr *)icmp_packet + ETHER_HDR_SIZE;
        struct icmp_hdr *icmph = (struct icmp_hdr *)icmp_packet + ETHER_HDR_SIZE + IP_HDR_SIZE(iph_temp);
        eh->ether_type = htons(ETH_P_IP);
        memcpy(eh->ether_shost, eh_temp->ether_shost, ETH_ALEN * sizeof(u8));
        memcpy(((char *)icmph) + 8, (char *)iph_temp, IP_HDR_SIZE(iph_temp) + 8);
        u32 daddr_temp = sip;
        u32 saddr_temp = ntohl(iph_temp->saddr);
        ip_init_hdr(iph, daddr_temp, saddr_temp, offset - ETHER_HDR_SIZE, IPPROTO_ICMP);
        memset(((char *)icmph) + 4, 0, 4);
        icmph->type = type;
        icmph->code = code;
        icmph->checksum = icmp_checksum(icmph, offset - ETHER_HDR_SIZE - IP_HDR_SIZE(iph_temp));
    }
    ip_send_packet(icmp_packet, offset);
}

```

代码 Ip\_forwarding.c

```

// forward the IP packet from the interface specified by longest_prefix_match,
// when forwarding the packet, you should check the TTL, update the checksum,
// determine the next hop to forward the packet, then send the packet by
// iface_send_packet_by_arp
void ip_forward_packet(u32 ip_dst, char *packet, int len)
{
    fprintf(stderr, "T000: forward ip packet.\n");

    struct ether_header *eh = (struct ether_header *)packet;
    struct iphdr *iph = packet_to_ip_hdr(packet);
    struct icmp_hdr *my_icmph = (struct icmp_hdr *)packet + ETHER_HDR_SIZE + IP_HDR_SIZE(iph);

    // deal with the packet
    u32 dst = ntohl(iph->daddr);
    rt_entry_t *my_entry = longest_prefix_match(dst);
    // failed
    if (!my_entry) {
        u32 dst = ntohl(iph->saddr);
        rt_entry_t *my_entry = longest_prefix_match(dst);
        u32 sip = my_entry->iface->ip;
        icmp_send_packet(packet, len, 3, 0, sip);
        free(packet);
        return;
    }
    // TTL -- 1 == 0
    iph->ttl -= 1;
    if (iph->ttl <= 0) {
        u32 dst = ntohl(iph->saddr);
        rt_entry_t *my_entry = longest_prefix_match(dst);
        u32 sip = my_entry->iface->ip;

        icmp_send_packet(packet, len, 11, 0, sip);
        free(packet);
        return;
    }
    iph->checksum = ip_checksum(iph);
    memcpy(eh->ether_shost, my_entry->iface->mac, ETH_ALEN);
    u32 next_hop = my_entry->ogw;
    if (!next_hop)
        next_hop = ntohl(iph->daddr);
    iface_send_packet_by_arp(my_entry->iface, next_hop, packet, len);
}

```

- 处理ip包

如果数据包是ICMP echo请求，目的IP地址等于iface的IP地址，则发送ICMP echo reply; 否则，转发数据包。



```

//handle ip packet
//
//If the packet is ICMP echo request and the destination IP address is equal to
//the IP address of the iface, send ICMP echo reply; otherwise, forward the
//packet.
void handle_ip_packet(iface_info_t *iface, char *packet, int len)
{
    struct ether_header *eh = (struct ether_header *)packet;
    struct iphdr *iph = packet_to_ip_hdr(packet);
    struct icmphdr *my_icmph = (struct icmphdr *) (packet + ETHER_HDR_SIZE + IP_HDR_SIZE(iph));

    if (my_icmph->type == 8 && ntohl(iph->daddr) == iface->ip) {
        u32 dst = ntohl(iph->saddr);
        rt_entry_t *my_entry = longest_prefix_match(dst);
        u32 sip = my_entry->iface->ip;

        icmp_send_packet(packet, len, 0, 0, sip);
        free(packet);
        return;
    }
    else {
        ip_forward_packet(iph->daddr, packet, len);
    }
}

```

## 代码 ip.c

- 在路由表中查找，找到具有相同和最长前缀的条目。输入地址按主机字节顺序排列

```

//lookup in the routing table, to find the entry with the same and longest prefix.
//the input address is in host byte order
rt_entry_t *longest_prefix_match(u32 dst)
{
    fprintf(stderr, "TODO: longest prefix match for the packet.\n");

    rt_entry_t *input = NULL;
    rt_entry_t *output = NULL;
    u32 number = 0;
    list_for_each_entry(input, &rttable, list) {
        u32 net = input->dest & input->mask;
        u32 net1 = input->mask & dst;
        if (net1 == net) {
            if (input->mask > number) {
                number = input->mask;
                output = input;
            }
        }
    }
    return output;
    //return NULL;
}

```

- 发送IP包

与ip\_forward\_packet不同，ip\_send\_packet发送路由器自身生成的数据包。此函数用于发送ICMP数据包。

```

//send IP packet
//
//Different from ip_forward_packet, ip_send_packet sends packet generated by
//router itself. This function is used to send ICMP packets.
void ip_send_packet(char *packet, int len)
{
    fprintf(stderr, "TODO: send ip packet.\n");

    struct iphdr *ip = packet_to_ip_hdr(packet);
    u32 dst = ntohl(ip->daddr);
    rt_entry_t *entry = longest_prefix_match(dst);
    if (!entry) {
        log(ERROR, "Could not find forwarding rule for IP (dst: 'IP_FMT') packet.",
            HOST_IP_FMT_STR(dst));
        free(packet);
        return;
    }

    u32 next_hop = entry->gw;
    if (!next_hop)
        next_hop = dst;

    struct ether_header *eh = (struct ether_header *)packet;
    eh->ether_type = htons(ETH_P_IP);
    struct iphdr *iph =
        (struct iphdr *) (packet + ETHER_HDR_SIZE);
    iph->saddr = htonl(entry->iface->ip);
    memcpy(eh->ether_shost, entry->iface->mac, ETH_ALEN * sizeof(u8));
    iface_send_packet_by_arp(entry->iface, next_hop, packet, len);
}

```

- 拓扑结构

```
1  from mininet.topo import Topo
2  from mininet.net import Mininet
3  from mininet.cli import CLI
4
5  class RouterTopo(Topo):
6      def build(self):
7          h1 = self.addHost('h1')
8          h2 = self.addHost('h2')
9          r1 = self.addHost('r1')
10         r2 = self.addHost('r2')
11
12         self.addLink(h1, r1)
13         self.addLink(h2, r2)
14         self.addLink(r1, r2)
15
16  if __name__ == '__main__':
17      topo = RouterTopo()
18      net = Mininet(topo = topo, controller = None)
19
20      h1, h2, r1, r2 = net.get('h1', 'h2', 'r1', 'r2')
21      h1.cmd('ifconfig h1-eth0 10.0.1.11/24')
22      h2.cmd('ifconfig h2-eth0 10.0.2.22/24')
23
24      h1.cmd('route add default gw 10.0.1.1')
25      h2.cmd('route add default gw 10.0.2.1')
26
27      for h in (h1, h2):
28          h.cmd('./scripts/disable_offloading.sh')
29          h.cmd('./scripts/disable_ipv6.sh')
30
31      r1.cmd('ifconfig r1-eth0 10.0.1.1/24')
32      r1.cmd('ifconfig r1-eth1 10.0.3.1/24')
33      r1.cmd('route add -
net 10.0.2.0 netmask 255.255.255.0 gw 10.0.3.2 dev r1-eth1')
34
35      r2.cmd('ifconfig r2-eth0 10.0.2.1/24')
36      r2.cmd('ifconfig r2-eth1 10.0.3.2/24')
37      r2.cmd('route add -
net 10.0.1.0 netmask 255.255.255.0 gw 10.0.3.1 dev r2-eth1')
38
39      r1.cmd('./scripts/disable_arp.sh')
40      r1.cmd('./scripts/disable_icmp.sh')
41      r1.cmd('./scripts/disable_ip_forward.sh')
42
43      r2.cmd('./scripts/disable_arp.sh')
44      r2.cmd('./scripts/disable_icmp.sh')
45      r2.cmd('./scripts/disable_ip_forward.sh')
```

```
46
47     net.start()
48     CLI(net)
49     net.stop()
```

## 五 实验结果

### 5.1 实验一

执行router\_topo.py文件，显示mininet下节点

```
root@ubuntu:~/networking/all_projects/cs/2-router# python router_topo.py
mininet> nodes
available nodes are:
h1 h2 h3 r1
```

显示组网信息

```
mininet> net
h1 h1-eth0:r1-eth0
h2 h2-eth0:r1-eth1
h3 h3-eth0:r1-eth2
r1 r1-eth0:h1-eth0 r1-eth1:h2-eth0 r1-eth2:h3-eth0
```

打开所有节点终端

```
mininet> xterm h1 h2 h3 r1
```

H1节点的ifconfig结果，显示节点的IP地址

```
root@ubuntu:~/networking/all_projects/cs/2-router# ifconfig
h1-eth0  Link encap:以太网 硬件地址 a2:96:21:eb:03:e2
          inet 地址:10.0.1.11 广播:10.0.1.255 掩码:255.255.255.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  点数:1
          接收数据包:12  错:0  弃:0  错:0
          发送数据包:1  错:0  弃:0  错:0
          接收:0  发送:0  队列:1000
          接收字节:936 (936.0 B)  发送字节:90 (90.0 B)

lo        Link encap:本地环回
          inet 地址:127.0.0.1 掩码:255.0.0.0
          UP LOOPBACK RUNNING  MTU:65536  点数:1
          接收数据包:0  错:0  弃:0  错:0
          发送数据包:0  错:0  弃:0  错:0
          接收:0  发送:0  队列:1000
          接收字节:0 (0.0 B)  发送字节:0 (0.0 B)

root@ubuntu:~/networking/all_projects/cs/2-router#
```

H2节点的ifconfig结果

```
"Node: h2"
root@ubuntu:~/networking/all_projects/cs/2-router# ifconfig
h2-eth0  Link encap:以太网  硬件地址 76:1c:58:8c:58:dc
         inet 地址:10.0.2.22  广播:10.0.2.255  掩码:255.255.255.0
         UP BROADCAST RUNNING MULTICAST  MTU:1500  点数:1
         接收数据包:12  发送:0  丢弃:0  接收:0
         发送数据包:1  发送:0  丢弃:0  发送:0
         接收:0  发送:0  列:度:1000
         接收字:936 (936.0 B)  发送字:90 (90.0 B)

lo       Link encap:本地回
         inet 地址:127.0.0.1  掩码:255.0.0.0
         UP LOOPBACK RUNNING  MTU:65536  点数:1
         接收数据包:0  发送:0  丢弃:0  接收:0
         发送数据包:0  发送:0  丢弃:0  发送:0
         接收:0  发送:0  列:度:1000
         接收字:0 (0.0 B)  发送字:0 (0.0 B)

root@ubuntu:~/networking/all_projects/cs/2-router#
```

H3节点的ifconfig结果

```
"Node: h3"
root@ubuntu:~/networking/all_projects/cs/2-router# ifconfig
h3-eth0  Link encap:以太网  硬件地址 ee:54:c4:c1:15:8f
         inet 地址:10.0.3.33  广播:10.0.3.255  掩码:255.255.255.0
         UP BROADCAST RUNNING MULTICAST  MTU:1500  点数:1
         接收数据包:12  发送:0  丢弃:0  接收:0
         发送数据包:1  发送:0  丢弃:0  发送:0
         接收:0  发送:0  列:度:1000
         接收字:936 (936.0 B)  发送字:90 (90.0 B)

lo       Link encap:本地回
         inet 地址:127.0.0.1  掩码:255.0.0.0
         UP LOOPBACK RUNNING  MTU:65536  点数:1
         接收数据包:0  发送:0  丢弃:0  接收:0
         发送数据包:0  发送:0  丢弃:0  发送:0
         接收:0  发送:0  列:度:1000
         接收字:0 (0.0 B)  发送字:0 (0.0 B)

root@ubuntu:~/networking/all_projects/cs/2-router#
```

R1节点的ifconfig结果，其包括3个网口eth0，eth1和eth2分别连接h1，h2，h3

```
root@ubuntu:~/networking/all_projects/cs/2-router# ifconfig
lo        Link encap:本地回环
          inet 地址:127.0.0.1 掩码:255.0.0.0
          inet6 地址: ::1/128 Scope:Host
          UP LOOPBACK RUNNING MTU:65536 点数:1
          接收数据包:0 丢弃:0 发送数据包:0 接收:0 发送:0
          接收字节:0 (0.0 B) 发送字节:0 (0.0 B)

r1-eth0    Link encap:以太网 硬件地址 72:08:68:a2:74:60
          inet 地址:10.0.1.1 广播:10.0.1.255 掩码:255.255.255.0
          inet6 地址: fe80::7008:68ff:fea2:7460/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST MTU:1500 点数:1
          接收数据包:1 丢弃:0 发送数据包:12 接收:0 发送:0
          接收字节:90 (90.0 B) 发送字节:936 (936.0 B)

r1-eth1    Link encap:以太网 硬件地址 5a:43:62:3e:65:fd
          inet 地址:10.0.2.1 广播:10.0.2.255 掩码:255.255.255.0
          inet6 地址: fe80::5843:62ff:fe3e:65fd/64 Scope:Link
          UP BROADCAST RUNNING MULTICAST MTU:1500 点数:1
          接收数据包:1 丢弃:0 发送数据包:0 接收:0 发送:0
```

在R1上执行router程序

```
root@ubuntu:~/networking/all_projects/cs/2-router# ./router
DEBUG: find the following interfaces: r1-eth0 r1-eth1 r1-eth2.
Routing table of 3 entries has been loaded.
```

在h1节点上ping自己显示的结果，表明能够ping通

```
root@ubuntu:~/networking/all_projects/cs/2-router# ping 10.0.1.11
PING 10.0.1.11 (10.0.1.11) 56(84) bytes of data.
64 bytes from 10.0.1.11: icmp_seq=1 ttl=64 time=0.030 ms
64 bytes from 10.0.1.11: icmp_seq=2 ttl=64 time=0.030 ms
64 bytes from 10.0.1.11: icmp_seq=3 ttl=64 time=0.036 ms
^C
--- 10.0.1.11 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2056ms
rtt min/avg/max/mdev = 0.030/0.032/0.036/0.003 ms
```

在h1上ping h2节点的IP地址，结果表明能够ping通，router程序执行正常

```
root@ubuntu:~/networking/all_projects/cs/2-router# ping 10.0.2.22
PING 10.0.2.22 (10.0.2.22) 56(84) bytes of data.
64 bytes from 10.0.2.22: icmp_seq=1 ttl=63 time=0.205 ms
64 bytes from 10.0.2.22: icmp_seq=2 ttl=63 time=0.121 ms
64 bytes from 10.0.2.22: icmp_seq=3 ttl=63 time=0.087 ms
64 bytes from 10.0.2.22: icmp_seq=4 ttl=63 time=0.130 ms
64 bytes from 10.0.2.22: icmp_seq=5 ttl=63 time=0.084 ms
64 bytes from 10.0.2.22: icmp_seq=6 ttl=63 time=0.082 ms
^C
--- 10.0.2.22 ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5107ms
rtt min/avg/max/mdev = 0.082/0.118/0.205/0.043 ms
```

在h1上ping h3的结果

```
root@ubuntu:~/networking/all_projects/cs/2-router# ping 10.0.3.33
PING 10.0.3.33 (10.0.3.33) 56(84) bytes of data.
64 bytes from 10.0.3.33: icmp_seq=1 ttl=63 time=0.182 ms
64 bytes from 10.0.3.33: icmp_seq=2 ttl=63 time=0.111 ms
^C
--- 10.0.3.33 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1003ms
rtt min/avg/max/mdev = 0.111/0.146/0.182/0.037 ms
```

在h1节点ping在网段中不存在的节点ip时，显示Destination Host Unreachable





```
"Node: h2"
root@ubuntu:~/networking/all_projects/cs/2-router# ifconfig
h2-eth0  Link encap:以太网 硬件地址 0e:15:84:d8:62:44
         inet 地址:10.0.2.22 广播:10.0.2.255 掩码:255.255.255.0
         UP BROADCAST RUNNING MULTICAST  MTU:1500  点数:1
         接收数据包:12  丢弃:0  接收错误:0  接收帧:0
         发送数据包:1  丢弃:0  发送错误:0  发送帧:0
         接收字节:976 (976.0 B)  发送字节:90 (90.0 B)

lo        Link encap:本地回环
         inet 地址:127.0.0.1 掩码:255.0.0.0
         UP LOOPBACK RUNNING  MTU:65536  点数:1
         接收数据包:0  丢弃:0  接收错误:0  接收帧:0
         发送数据包:0  丢弃:0  发送错误:0  发送帧:0
         接收字节:0 (0.0 B)  发送字节:0 (0.0 B)

root@ubuntu:~/networking/all_projects/cs/2-router#
```

R1节点显示如下

```
"Node: r1"
接收数据包:0  丢弃:0  接收错误:0  接收帧:0
发送数据包:0  丢弃:0  发送错误:0  发送帧:0
接收字节:0 (0.0 B)  发送字节:0 (0.0 B)

r1-eth0  Link encap:以太网 硬件地址 1a:52:2f:82:24:d1
         inet 地址:10.0.1.1 广播:10.0.1.255 掩码:255.255.255.0
         inet6 地址: fe80::1852:2fff:fe82:24d1/64 Scope:Link
         UP BROADCAST RUNNING MULTICAST  MTU:1500  点数:1
         接收数据包:1  丢弃:0  接收错误:0  接收帧:0
         发送数据包:10  丢弃:0  发送错误:0  发送帧:0
         接收字节:90 (90.0 B)  发送字节:796 (796.0 B)

r1-eth1  Link encap:以太网 硬件地址 82:d8:1f:99:76:07
         inet 地址:10.0.3.1 广播:10.0.3.255 掩码:255.255.255.0
         inet6 地址: fe80::80d8:1fff:fe99:7607/64 Scope:Link
         UP BROADCAST RUNNING MULTICAST  MTU:1500  点数:1
         接收数据包:12  丢弃:0  接收错误:0  接收帧:0
         发送数据包:10  丢弃:0  发送错误:0  发送帧:0
         接收字节:976 (976.0 B)  发送字节:796 (796.0 B)

root@ubuntu:~/networking/all_projects/cs/2-router#
```

R2节点显示如下



```
*****
TODO: forward ip packet.
TODO: longest prefix match for the packet.
TODO: lookup ip address in arp cache.
arp cache is:
10.0.3.1      7a:96:f6:05:24:88      1558492325      1
10.0.2.22     6e:c2:31:8b:ef:ee      1558492326      1
*****
```